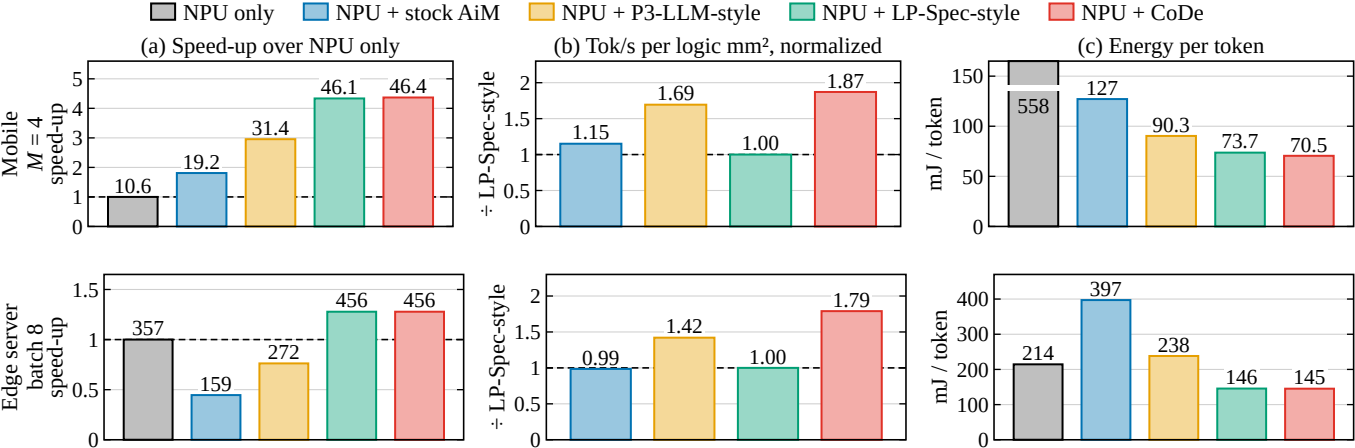
NPU only
NPU + stock AiM
NPU + P3-LLM-style
NPU + LP-Spec-style
NPU + CoDe
(a) Speed-up over NPU only
(b) Tok/s per logic mm², normalized
(c) Energy per token
Mobile
M = 4
speed-up
10.6
19.2
31.4
46.1
46.4
÷ LP-Spec-style
1.15
1.69
1.00
1.87
mJ / token
558
127
90.3
73.7
70.5
Edge server
batch 8
speed-up
357
159
272
456
456
÷ LP-Spec-style
0.99
1.42
1.00
1.79
mJ / token
214
397
238
146
145