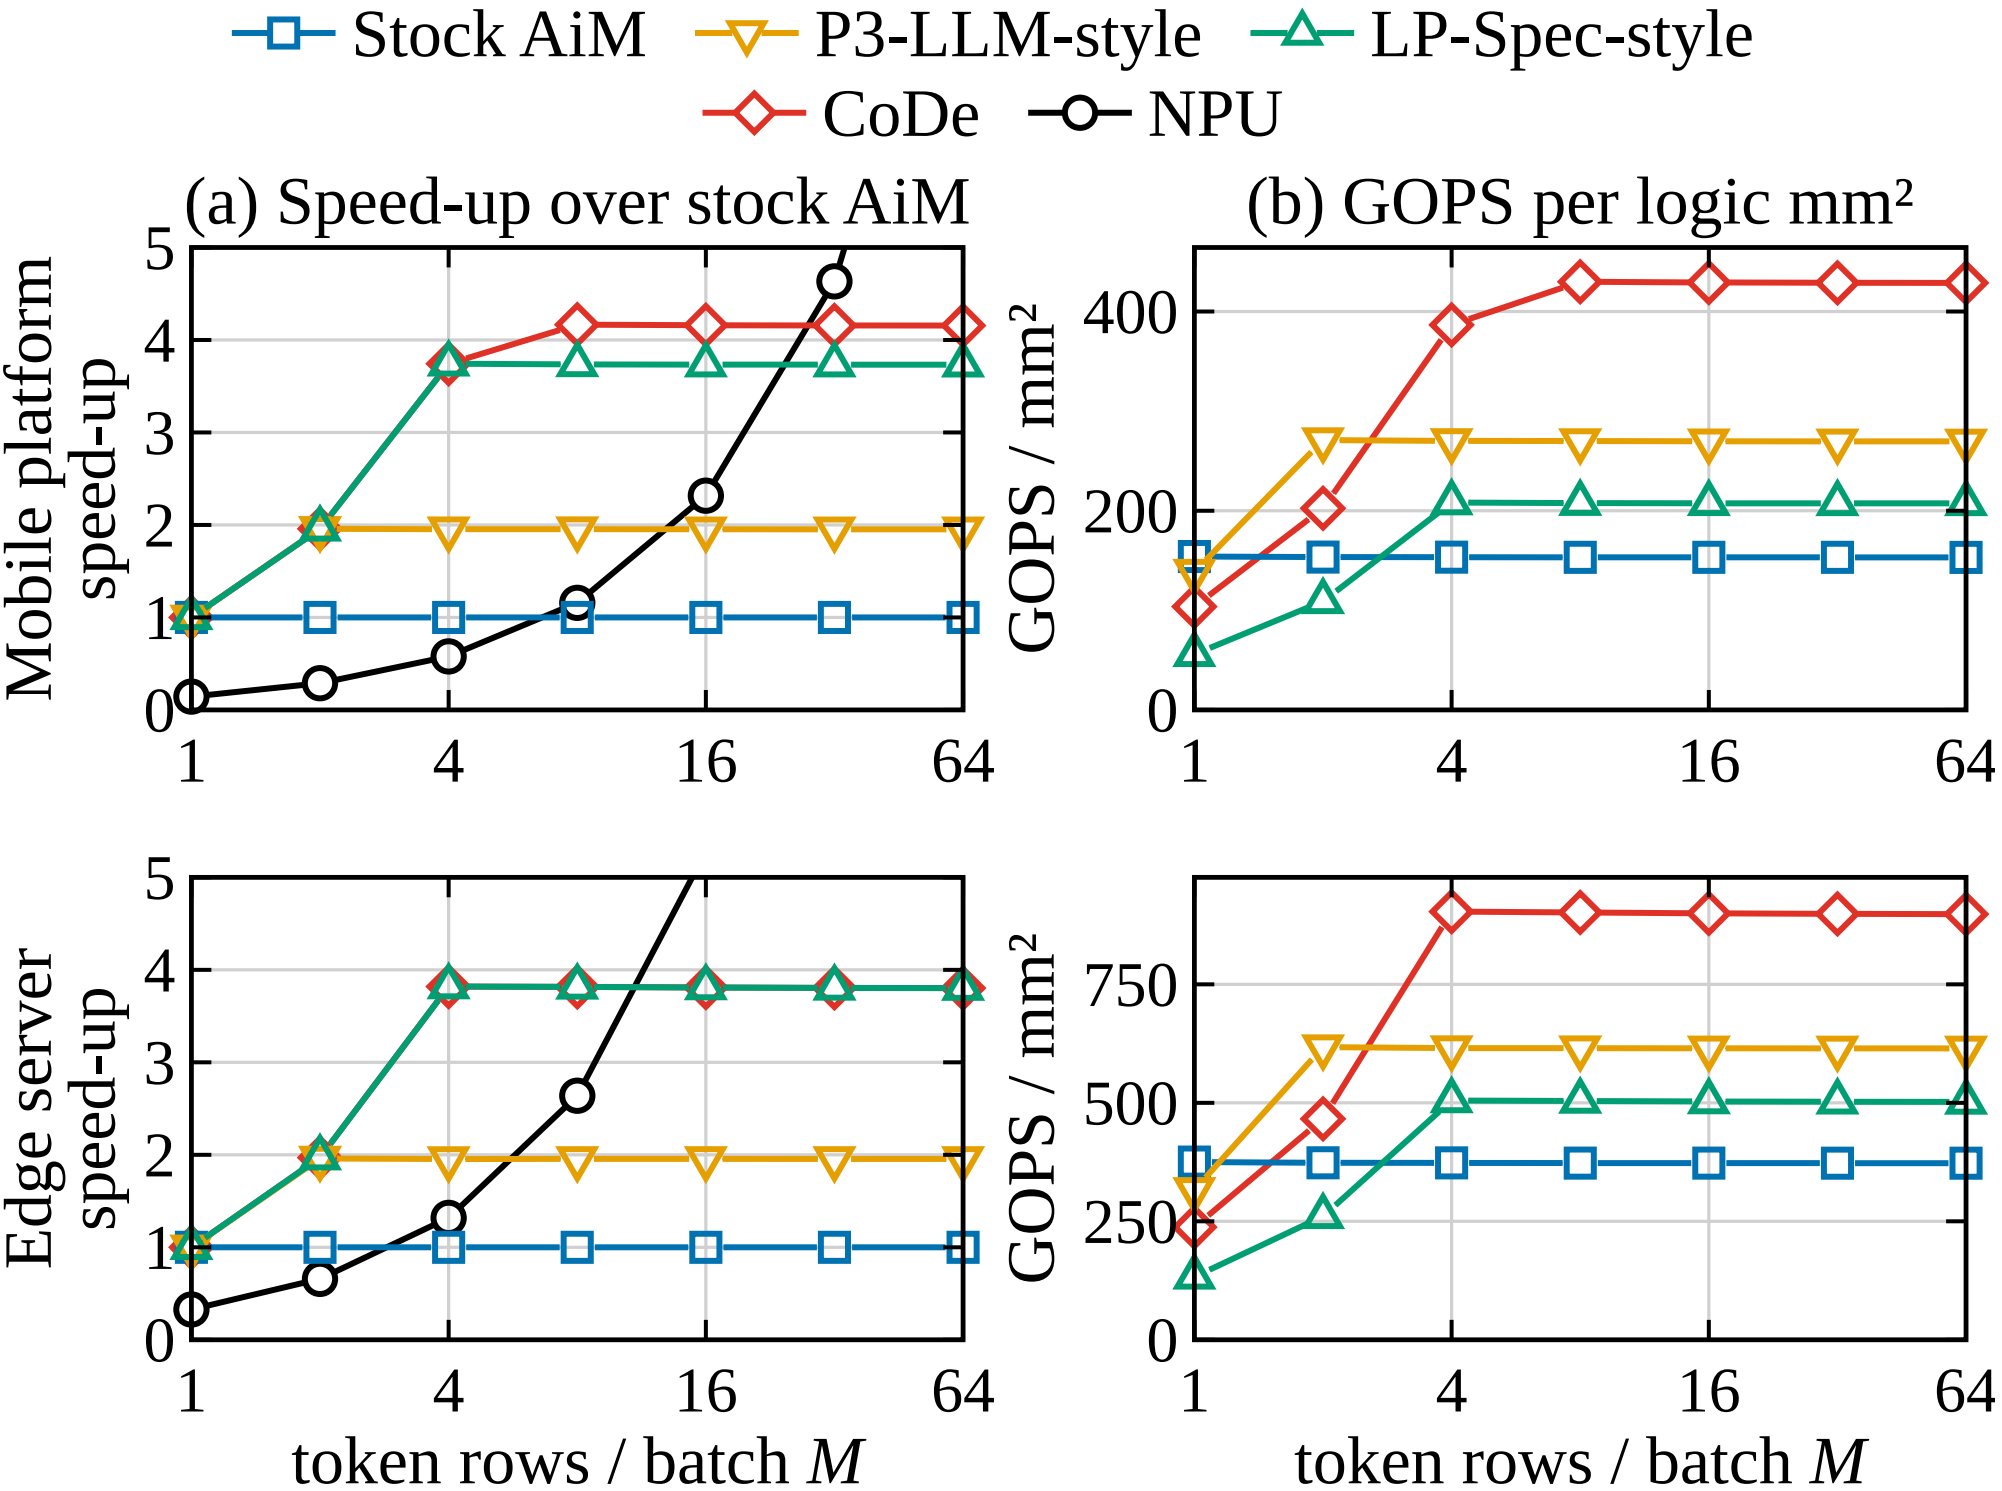
Stock AiM
P3-LLM-style
LP-Spec-style
CoDe
NPU
(a) Speed-up over stock AiM
(b) GOPS per logic mm²
Mobile platform speed-up
Edge server speed-up
GOPS / mm²
token rows / batch M
0
1
2
3
4
5
200
400
250
500
750
1
4
16
64