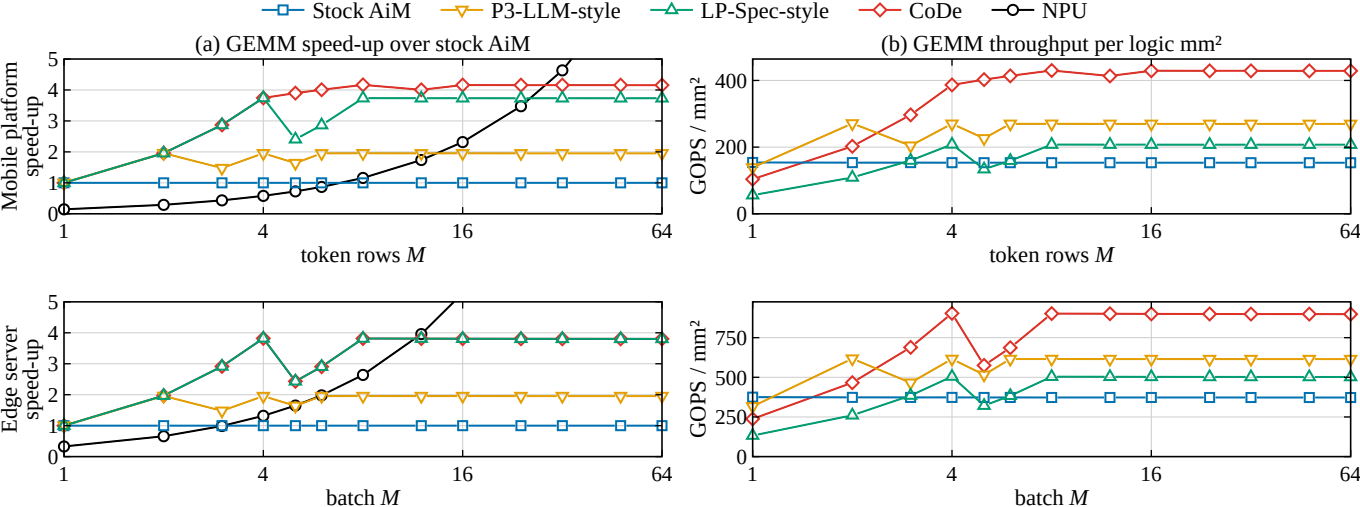
Stock AiM
P3-LLM-style
LP-Spec-style
CoDe
NPU
(a) GEMM speed-up over stock AiM
(b) GEMM throughput per logic mm²
Mobile platform speed-up
GOPS / mm²
token rows M
Edge server speed-up
batch M
0
1
2
3
4
5
200
400
500
750
1
4
16
64