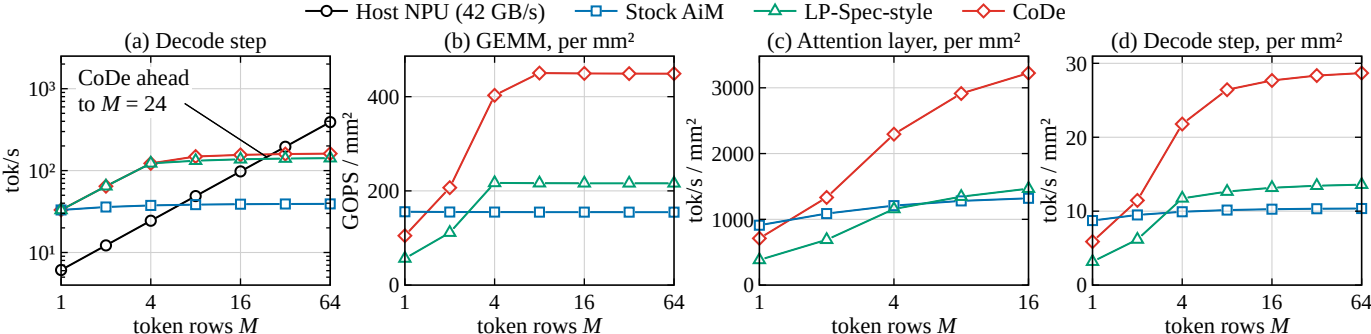
Host NPU (42 GB/s)
Stock AiM
LP-Spec-style
CoDe
(a) Decode step
CoDe ahead
to M = 24
tok/s
token rows M
(b) GEMM, per mm²
GOPS / mm²
token rows M
(c) Attention layer, per mm²
tok/s / mm²
token rows M
(d) Decode step, per mm²
tok/s / mm²
token rows M